

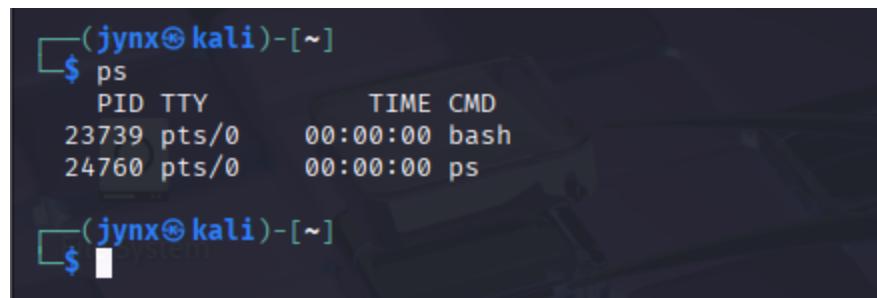
CHAPTER 6 - PROCESS MANAGEMENT

A **process** is essentially a program/thread that's running and using resources [CPU, GPU, RAM etc.], an efficient Linux administrator, digital forensic investigator or **hackers** need to have a particularly relevant and profound understanding of processes and how they work.

▼ Viewing Processes.

The '**ps**' command is used to observe and analyze which background processes are currently running on the system.

example:



```
(jynx㉿kali)-[~]
$ ps
  PID TTY      TIME CMD
 23739 pts/0    00:00:00 bash
 24760 pts/0    00:00:00 ps

(jynx㉿kali)-[~]
$
```

The Linux kernel assigns unique **Process IDs** (PIDs) sequentially to each process as it's created. When managing processes, PIDs are more important

than process names for identification.

Running

`ps` alone provides minimal information—only processes started by the current user on the current terminal (typically just the bash shell and the `ps` command itself).

For complete system visibility, use; `ps aux` - this command displays all processes running on the system for all users, including:

- System background processes
- Processes from other users
- Detailed process information

Note: The options `aux` are used ***without*** a dash (-) and must be lowercase, as Linux is case-sensitive.

```
(jynx㉿kali)-[~]
$ ps aux
USER      PID %CPU %MEM    VSZ   RSS TTY      STAT START  TIME COMMAND
root      1  0.0  0.3 23568 13980 ?        Ss  08:59  0:00 /sbin/init splash
root      2  0.0  0.0     0    0 ?        S   08:59  0:00 [kthreadd]
root      3  0.0  0.0     0    0 ?        S   08:59  0:00 [pool_workqueue_release]
root      4  0.0  0.0     0    0 ?        I<  08:59  0:00 [kworker/R-kvfree_rcu_reclaim]
root      5  0.0  0.0     0    0 ?        I<  08:59  0:00 [kworker/R-rcu_gp]
root      6  0.0  0.0     0    0 ?        I<  08:59  0:00 [kworker/R-sync_wq]
root      7  0.0  0.0     0    0 ?        I<  08:59  0:00 [kworker/R-slub_flushwq]
root      8  0.0  0.0     0    0 ?        I<  08:59  0:00 [kworker/R-netns]
root     10  0.0  0.0     0    0 ?        I<  08:59  0:00 [kworker/0:0H-kblockd]
root     12  0.0  0.0     0    0 ?        I   08:59  0:00 [kworker/u128:0-ipv6_addrconf]
root     13  0.0  0.0     0    0 ?        I<  08:59  0:00 [kworker/R-mm_percpu_wq]
root     14  0.0  0.0     0    0 ?        I   08:59  0:00 [rcu_tasks_kthread]
root     15  0.0  0.0     0    0 ?        I   08:59  0:00 [rcu_tasks_rude_kthread]
root     16  0.0  0.0     0    0 ?        I   08:59  0:00 [rcu_tasks_trace_kthread]
root     17  0.0  0.0     0    0 ?        S   08:59  0:00 [ksoftirqd/0]
root     18  0.0  0.0     0    0 ?        I   08:59  0:00 [rcu_prempt]
root     19  0.0  0.0     0    0 ?        S   08:59  0:00 [rcu_exp_par_gp_kthread_worker/1]
root     20  0.0  0.0     0    0 ?        S   08:59  0:00 [rcu_exp_gp_kthread_worker]
root     21  0.0  0.0     0    0 ?        S   08:59  0:00 [migration/0]
root     22  0.0  0.0     0    0 ?        S   08:59  0:00 [idle_inject/0]
root     23  0.0  0.0     0    0 ?        S   08:59  0:00 [cpuhp/0]
root     24  0.0  0.0     0    0 ?        S   08:59  0:00 [cpuhp/1]
root     25  0.0  0.0     0    0 ?        S   08:59  0:00 [idle_inject/1]
root     26  0.0  0.0     0    0 ?        S   08:59  0:00 [migration/1]
root     27  0.0  0.0     0    0 ?        S   08:59  0:00 [ksoftirqd/1]
root     29  0.0  0.0     0    0 ?        I<  08:59  0:00 [kworker/1:0H-events_highpri]
root     30  0.0  0.0     0    0 ?        S   08:59  0:00 [cpuhp/2]
root     31  0.0  0.0     0    0 ?        S   08:59  0:00 [idle_inject/2]
root     32  0.0  0.0     0    0 ?        S   08:59  0:00 [migration/2]
root     33  0.0  0.0     0    0 ?        S   08:59  0:00 [ksoftirqd/2]
root     35  0.0  0.0     0    0 ?        I<  08:59  0:00 [kworker/2:0H-events_highpri]
root     36  0.0  0.0     0    0 ?        S   08:59  0:00 [cpuhp/3]
root     37  0.0  0.0     0    0 ?        S   08:59  0:00 [idle_inject/3]
root     38  0.0  0.0     0    0 ?        S   08:59  0:00 [migration/3]
root     39  0.0  0.0     0    0 ?        S   08:59  0:00 [ksoftirqd/3]
root     41  0.0  0.0     0    0 ?        I<  08:59  0:00 [kworker/3:0H-events_highpri]
root     44  0.0  0.0     0    0 ?        I   08:59  0:00 [kworker/u131:0-flush-8:0]
```

COLUMNS (understanding) in the output:

1. **USER** - The user who invoked the process.
2. **PID** - The process ID [unique].
3. **%CPU** - Percentage of CPU being utilized by the process.
4. **%MEM** - Percentage of memory being utilized by the process.
5. **COMMAND** - The name of the command that started the process.

→ To perform any manipulation commands on a particular process, we must specify its **PID**.

▼ Filtering by Process Name.

The **grep** command can be used to filter out processes on the basis of name(s).

example:

```
(jynx㉿kali)-[~]
└─$ ps aux | grep msfconsole
jynx      32710  0.0  0.0    6528   2240 pts/0    S+   10:07   0:00 grep --color=auto msfconsole
```

▼ 'top' command.

By default, the **ps** command displays processes in order they were started (chronologically), essentially by PIDs. If you wish to see the processes consuming most system resources use 'top' command.

```
> top
```

example:

File Actions Edit View Help												
top - 10:11:45 up 1:12, 1 user, load average: 0.00, 0.00, 0.00												
Tasks: 227 total, 1 running, 226 sleeping, 0 stopped, 0 zombie												
%Cpu(s): 3.2 us, 0.6 sy, 0.0 ni, 96.1 id, 0.0 wa, 0.0 hi, 0.2 si, 0.0 st												
MiB Mem : 3912.7 total, 2465.8 free, 917.9 used, 767.0 buff/cache												
MiB Swap: 953.7 total, 953.7 free, 0.0 used. 2994.8 avail Mem												
PID	USER	PR	NI	VIRT	RES	SHR	S	%CPU	%MEM	TIME+	COMMAND	
1018	root	20	0	448856	160920	69428	S	10.0	4.0	0:28.60	Xorg	
23703	jynx	20	0	733204	70104	53236	S	1.0	1.7	0:01.71	qterminal	
518	root	20	0	0	0	0	I	0.3	0.0	0:01.04	kworker/2:2-mm_percpu_wq	
1300	jynx	20	0	168744	7600	6960	S	0.3	0.2	0:00.11	at-spi2-registr	
1320	jynx	20	0	1189344	127524	85084	S	0.3	3.2	0:09.94	xfwm4	
1379	jynx	20	0	311504	65336	22936	S	0.3	1.6	0:11.64	wrapper-2.0	
1383	jynx	20	0	425628	37560	29768	S	0.3	0.9	0:00.13	wrapper-2.0	
1472	jynx	20	0	153356	42408	31508	S	0.3	1.1	0:10.35	vmtoolsd	
34534	jynx	20	0	10488	5796	3620	R	0.3	0.1	0:00.03	top	
1	root	20	0	23568	13980	10544	S	0.0	0.3	0:00.96	systemd	
2	root	20	0	0	0	0	S	0.0	0.0	0:00.01	kthreadd	
3	root	20	0	0	0	0	S	0.0	0.0	0:00.00	pool_workqueue_release	
4	root	0	-20	0	0	0	I	0.0	0.0	0:00.00	kworker/R-kvfree_rcu_reclaim	
5	root	0	-20	0	0	0	I	0.0	0.0	0:00.00	kworker/R-rcu_gp	

→ While you have top running, pressing the `H` or `?` key will bring up a list of interactive commands, and pressing `Q` will quit top.

▼ Managing Processes.

The `nice` command is used to change the priority of a process to the kernel.

The values for nice range from `-20` to `+19`, with zero being the default value.

`-20` - Most likely to receive priority

`0` - Default priority

`-19` - Least likely to receive priority

→ You can also alter the priority using `renice` command.

→ The `nice` command requires to increment the value, on the other hand the `renice` command wants an absolute value.

example:

```
| nice -n -15 /bin/slowproces123
```

This command above, would increment by the `nice` value by `15` and prioritize process- slowprocess123.

```
| nice -n 10 /bin/slowproces123
```

This command above, would decrement by the `nice` value by 10 and de-prioritize the process- slowprocess123.

▼ Killing Processes.

Processes can become problematic when they:

- Consume excessive system resources
- Display abnormal behavior
- Become unresponsive or freeze

Such problematic processes waste valuable system resources that could be allocated to functional processes, they are called **zombie processes**.

When you identify a problematic process, use the `kill` command to terminate it. The kill command offers multiple termination methods, each identified by a specific signal number that determines how the process should be stopped.

Syntax: `kill -signal PID`, signal switch is optional.

→ Common KILL Signals:

Signal Name	Number	Description
SIGHUP	1	Hangup Signal - Terminates the process and immediately restarts it using the same Process ID. Commonly used to reload configuration files.
SIGINT	2	Interrupt Signal - A gentle termination request that allows the process to clean up before exiting. Similar to pressing Ctrl+C. May not work if the process is unresponsive.
SIGQUIT	3	Quit Signal - Forces process termination and creates a core dump file in the current directory. The core dump contains the process's memory state for debugging purposes.
SIGTERM	15	Termination Signal - The standard and default kill signal. Politely requests the process to terminate, allowing it to perform cleanup operations before shutting down.
SIGKILL	9	Kill Signal - The most forceful termination method. Immediately destroys the process without allowing

cleanup. Use as a last resort when other signals fail.

EXAMPLE:

- To restart a process with the HUP signal for a process with PID 6996, enter the -1 option with kill:

`kill -1 6996`

▼ Running Processes in the Background.

In Linux systems, all user interactions occur through a shell environment, regardless of whether you're using the command line interface or graphical desktop. Even GUI applications ultimately execute their commands through the underlying shell.

When you run a command, the shell operates synchronously by default - it waits for the current command to finish executing before displaying the next command prompt and accepting new input.

Sometimes you need to run a process while continuing to use the terminal for other tasks. For example, if you start a file download with `wget` :

```
 wget https://example.com/largefile.zip
```

The terminal becomes occupied with the download process, showing progress updates but preventing you from entering new commands until the download completes.

Instead of opening multiple terminals or waiting for the process to finish, you can run processes in the background. Background processes continue running independently while freeing up the terminal for other commands.

To run a command in the background, append an ampersand (&) to the end:

```
 wget https://example.com/largefile.zip &
```

Now the download runs in the background, and you immediately get a new command prompt to execute other tasks while the file continues downloading.

Benefits

This approach is particularly useful for:

- Long-running downloads or file transfers
- System monitoring tools
- Any process that takes significant time to complete
- Maximizing terminal efficiency without opening multiple windows

If you want to move a process running in the background to the *foreground*, you can use the `fg` (foreground) command. The `fg` command requires the PID of the process you want to return to the foreground.

| `fg 6996`

▼ Scheduling Processes

System administrators and developers frequently need to automate tasks at specific times. Common scenarios include:

- Running system backups during off-peak hours
- Performing regular maintenance tasks
- Executing monitoring scripts at intervals
- Automating data processing jobs

Scheduling Tools

Linux provides two primary scheduling mechanisms:

at Command

Used for **one-time** scheduled tasks. The `at` daemon runs commands at a specified future time.

cron/crond

Better suited for **recurring** tasks that need to run daily, weekly, or monthly.

The `at` command accepts various time formats for flexible scheduling:

Time Format	Meaning
at 2:30pm	Run at 2:30 PM today
at 2:30pm July 15	Run at 2:30 PM on July 15
at midnight	Run at midnight today
at midnight July 15	Run at midnight on July 15
at tomorrow	Run tomorrow at current time
at now + 30 minutes	Run in 30 minutes
at now + 2 hours	Run in 2 hours
at now + 3 days	Run in 3 days
at now + 2 weeks	Run in 2 weeks
at 2:30pm 07/15/2025	Run at 2:30 PM on July 15, 2025

Example Usage

To schedule a backup script to run at 3:00 AM:

```
$ at 3:00am
at> /home/user/backup_script.sh
at> <Ctrl+D>
```

The `at` command enters interactive mode where you specify the commands to execute, then press Ctrl+D to save the scheduled task.